

AKSHAY KHULE

Principal Architect | Full-Stack Development · Distributed Systems · Data Platforms · Generative AI

Mobile: 9326768815 | Email: Khule.akshay@gmail.com

[LinkedIn](#) | [GitHub](#) | [Cloud Certifications](#) | [Distributed Systems Certifications](#)



PROFESSIONAL SUMMARY

Principal Architect | Making complex distributed systems simple.

- Global technical leadership across Software Engineering, Product Development, and Systems Architecture — architecting and shipping distributed systems, cloud-native microservices, petabyte-scale data platforms, and GenAI solutions across Banking, Pharma, AIOps, Manufacturing, and Cybersecurity domains, driving end-to-end High-Level and Low-Level Design for scalable, resilient architectures.
- Architecture leadership across multiple product teams (led teams of 30+); delivered winning RFP responses for new business; accountable for technical design, product integration, and compliance with architecture and design principles.
- Hands-on product architecture and development of microservices and serverless applications - Java/J2EE, Spring Boot, Spring Cloud, Spring Security, OAuth/KeyCloak, JPA, REST, WebSockets, RDBMS with database schema design, and message queues -engineered with configuration management, service discovery and registration, resiliency patterns, and distributed log tracing, containerized with Docker and deployed on Kubernetes across AWS and Azure. Developed UX/UI with basic JavaScript and React knowledge.
- Expert in end-to-end Big Data, Data Lake, and Data Platform architecture - streaming/batch ingestion, storage, processing, governance, and MDM - using PySpark (Databricks), Pandas, Kafka, Delta Lake, Azure Synapse, Data warehouse, Azure Data Factory, and Microsoft Fabric, delivering petabyte-scale, GDPR-compliant pipelines with zero-data-loss guarantees.
- Hired, mentored, and guided Application Development and Data Platform teams (architects, lead developers, engineers). Defined architecture guardrails, engineering standards, and reference implementations adopted organization wide.
- Generative AI: Architected the multi-agent foundation for [Unisys Intelligence Accelerator](#) - FastAPI microservices exposing agent capabilities via REST, with a custom AWS agent harness running on AWS Bedrock AgentCore and Microsoft Foundry, grounded through managed Knowledge Bases (RAG) - automating IT operations, data tasks for large enterprises.

PUBLICATIONS & THOUGHT LEADERSHIP

- **White paper:** GenAI-Augmented Polyglot Persistence for Microservices - Toward Self-Adaptive & Explainable Data Fabrics. <https://lnkd.in/d/VeegA4H>
- **Newsletter:** The Engineering Log by Akshay - ongoing LinkedIn publication on distributed systems, data engineering, and GenAI. [Read on LinkedIn](#)

CERTIFICATIONS

AWS	Solutions Architect Professional Data Analytics Specialty Security Specialty
Databricks	Databricks Certified Professional Spark Certified Associate Developer
Confluent (Kafka)	Confluent Certified Developer for Apache Kafka
Microsoft Azure	Microsoft Certified: Azure Fundamentals
DAMA	Certified Data Management Professional (CDMP)
IIT Roorkee	Generative AI Engineering Program (completed)
Anthropic	Claude 101 - Certificate of Completion

TECHNICAL SKILLS

Languages	Java, Python, Scala, SQL, PL/SQL, Java Script, HTML
Big Data & Analytics	Kafka, PySpark, Pandas, Databricks, Delta Lake, Hive, HDFS, Hadoop, Medallion Lakehouse, Azure Synapse Analytics, Azure Data Factory, Microsoft Fabric (OneLake), AWS Glue, Redis, Elasticsearch, Data Governance, Microsoft Purview, S3, Azure Storage account, Vector Database, MDM
Cloud & Containers	AWS, Azure, Kubernetes (AKS), Docker, Helm, Microsoft Sentinel, Databricks Asset Bundles, GitHub Actions, Logic Apps, Azure Functions, AWS Lambda, Synapse Dedicated pool, Grafana, API Gateway
Frameworks & Tools	Microservices, FastAPI, Jupyter, Spring Framework, Spring Security, Flask, Hibernate ORM, EJB, JMS, MDB, REST, OAuth, KeyCloak, React, JWT, Jenkins, Maven, Git, SonarQube, SQLAlchemy
Generative AI	LLMs, Agentic AI, Multi-Agent Orchestration (Agno), RAG, AI Guardrails, AWS Bedrock, Microsoft Foundry, Azure AI Search, Claude Code, Hugging Face Transformers, LLM training/fine-tuning

PROFESSIONAL EXPERIENCE

Principal Architect — Unisys Corporation

May 2021 – Present

Architect (2021 – 2024); Principal Architect (2024 – Present)

Unisys CloudForte - Enterprise Monitoring & Security Product (Unisys Stealth Family)

Observability product in the CSPM/cybersecurity portfolio - monitors enterprise infrastructure, applications in real time, providing the telemetry foundation for anomaly detection, security monitoring, and threat visibility.

URL: marketplace.microsoft.com — CloudForte for Azure

- Designed the overall product architecture, cloud platform setup, and technology choices for CloudForte. Built the event-processing layer as microservices on Azure Kubernetes Service (AKS) - consuming ~5M events/day safely (no duplicates, automatic retries, dead-letter queues for failed messages) and triggering alerts or automated fixes based on risk. Extended security events into Microsoft Sentinel for SIEM correlation and SOAR automation, with per-tenant data isolation.
- Designed an event-driven GenAI incident-triage agent - microservices consumes incidents and invokes an Azure AI Foundry agent configured with runbook and incident-history knowledge sources- recommending fixes, and estimating impact, with the AI only recommending while actual fixes execute through approved, tested runbooks. Achieved 30% MTTR reduction.

Banking Data Platform Modernization - Financial Services Client

- Designed the migration strategy after Microsoft deprecated the Export to Data Lake service - built a data-ingestion solution pulling Dynamics 365 F&O data from Microsoft Fabric (Fabric Link/OneLake) into an Azure Synapse dedicated SQL pool, with zero disruption to downstream financial reporting across ~100+ pipelines.
- Developed a PySpark-based solution in Azure Databricks (Delta Lake) to process Dynamics 365 F&O financial data, migrating ledger and financial-transaction processing into scalable, version-controlled Spark workloads - cutting pipeline runtime from ~3 hrs to 20 mins. Implemented Databricks Asset Bundles (DABs) with GitHub Actions CI/CD to declaratively define, promote, and deployments across dev, staging, and prod environments. Implemented Databricks Lakehouse Federation.

Architect — Cognizant (Client: Novartis)

July 2018 – May 2021

Adverse Event resolution product for Novartis — analyses adverse events and predicts resolutions.

URL: <https://ae-service.novartis.net/pae-coreui>

- Architected and delivered a full-stack, event-driven system — Salesforce Social Studio Stream, microservices, AWS Lambda, API Gateway, and SQS over REST — driving cloud adoption that saved nearly \$1M annually.
- Led the Strategic Data Platform build-out, moving the stack from synchronous to asynchronous (SQS + Lambda) and designing an S3 data lake consolidating RDBMS and file-system sources, solving critical data-latency problems for the data-ingestion product - while driving Data Security, HIPAA, and GDPR compliance initiatives. Contributed hands-on across the full stack - UI, APIs, backend and data pipelines.

Technical Manager — Capgemini (Client: Morgan Stanley)

Feb 2017 – June 2018

- Owned end-to-end product development - architecture, design, development, release, and maintenance of Spring Boot microservices: defined service boundaries and API contracts, implemented REST controllers, service and DAO layers, and designed entities and the underlying data model, and full stack development using React JS, Spring boot.

Technical Lead — Persistent Systems (Client: Thermo Fisher/Proofpoint)

Sept 2014 – Jan 2017

Thermo Fisher- Ion Reporter: NGS product automating sequencing, analysis, and reporting. URL:

<https://ionreporter.thermofisher.com>

- Led backend development of REST APIs and Spring/Hibernate persistence layers for the sequencing analysis workflow
Proofpoint — Enterprise email security and compliance platform.
- Built high-volume message-processing and policy-management services for threat-detection and DLP/archiving modules, with idempotent processing, RBAC, and audit trails for regulated customers.

System Analyst — Worldline (Client: BNP Paribas)

Apr 2013 – Aug 2014

Worldline SEPA solutions - an all-in-one payment ecosystem that automates, secures, and streamlines Single Euro Payments Area (SEPA) transactions across Europe.

- Built a reusable Spring/Spring REST transaction engine and Hibernate ORM data layers for liquidity and transaction management - optimizing SEPA direct-debit and credit-transfer processing. Also contributed to front-end UI development.

Consultant — Polaris (Client: JPMorgan)

Mar 2011 – Mar 2013

TLIP banking application for JP Morgan - integration of the TIDE, LIQUIDITY, and PORTAL modules. URL: intellectdesign.com

- Designed core subsystem boundaries, application interfaces, and multi-tier enterprise runtime platforms using EJB and JMS; collaborated with QA on complex test analysis, defect resolution, and performance validation.

Consultant — HAL (in-contract)

Jan 2009 – Mar 2011

- Building an in-house inventory management and supply chain solution with batch application support for an ERP.

KEY ENGINEERING HIGHLIGHTS

Microservices & Backend

- Backend Stack Ownership: Took complete ownership of the backend stack - from development through cloud deployment - for 30+ microservices serving enterprise applications across multiple organizations; ensured resiliency via circuit breakers, bounded retries with exponential backoff, dead-letter queues, zero-downtime deployments on Kubernetes (AKS/EKS).
- Event-Driven Consistency: Built SQS-based microservices using the transactional outbox pattern - saving business data and the outgoing event in one database transaction before publishing

Distributed Systems Design

- High-Throughput Partitioning: diagnosed and fixed hot-partition skew in high-velocity Kafka pipelines by redesigning topic partition keys (entity-level composite keys).
- Fault-Tolerant Design: designed cloud data-ingestion services for partition tolerance - idempotent, retry-safe operations with exponential backoff and jitter, plus dead-letter queues.

Massive-Scale Data Ingestion, Engineering and Governance

- Petabyte-Scale Ingestion: architected a high-velocity ingestion engine on Spark handling petabytes of records.
- Low-Latency Streaming: implemented event-driven Kafka pipelines with fault-tolerant Java consumer-group services — cutting data availability to minutes. Utilize Azure Purview for Data Governance and MDM.

DevOps & CI/CD

- Infrastructure as Code: architected Databricks Asset Bundles to declaratively define, version, and promote data workloads across environments - integrated with secure GitHub Actions pipelines on VNet-isolated self-hosted runners with strict network boundaries, environment-specific variable injection. Used ADF arm template for Pipeline deployment.
- Containerized Orchestration: containerized distributed applications with Docker and orchestrated fault-tolerant deployments on managed Kubernetes (AKS), leveraging auto-scaling, rolling updates, and self-healing.

Generative AI

- Designed and Develop Agents on Azure AI Foundry/Amazon Bedrock, coordinating specialized retrieval, reasoning, and execution agents to autonomously handle complex operational workflows.

AWARDS & RECOGNITION

- **Unisys All-Hands Recognition:** recognized four times at company-wide All-Hands meetings for high-impact work across distributed systems and data engineering initiatives.
- **“You Rock” Award Nomination — Persistent:** nominated by peers and leadership for consistent technical excellence.

EDUCATION

- Bachelor of Engineering (B.E.), Computer Science — V.Y.W.S. College of Engineering, Badnera, Amravati University (1st Division).
- Generative AI Engineering Program — IIT Roorkee.